

# **INTRODUCTION TO DEEP LEARNING FOR COMPUTER VISION**

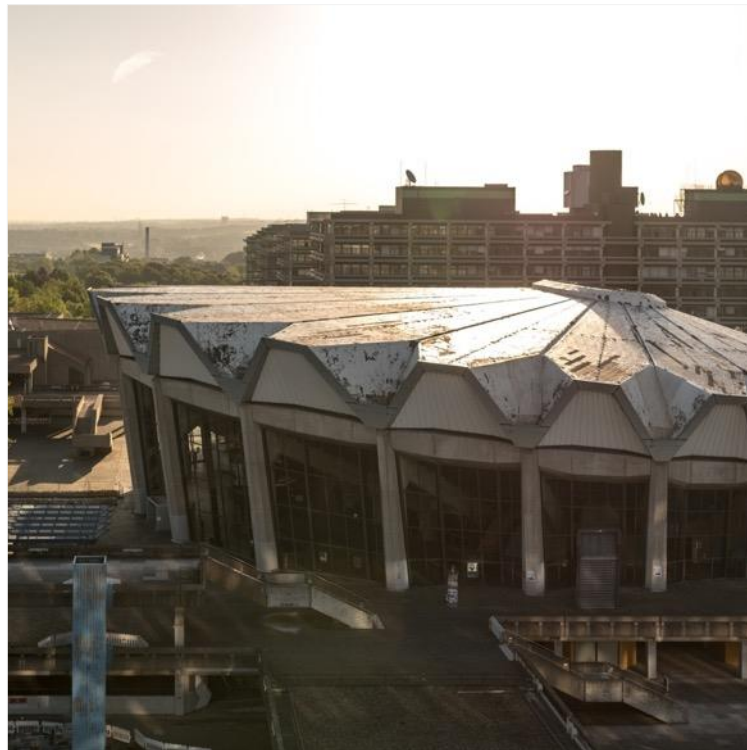
## **DAY 1 – BASICS**

SEBASTIAN HOUBEN

# Schedule



## Today

- Computer Vision and Deep Learning
- Image Classification
- Representation of images in Python
- Feature extraction
- Evaluating an image classifier
- Convolution
- German Traffic Sign Recognition Benchmark




# Computer Vision

- Programs that process images as input
- Gain understanding of images or video
- Mimic performance of human visual system
- Typical tasks
  - Object detection
  - Object segmentation
  - Image registration
  - Pose estimation
  - Face recognition
  - Egomotion
  - Optical Flow



# Computer Vision


- Programs that process images as input
- Gain understanding of images or video
- Mimic performance of human visual system
- Typical tasks
  - Object detection
  - Object segmentation
  - Image registration
  - Pose estimation
  - Face recognition
  - Egomotion
  - Optical Flow







# Computer Vision

- Programs that process images as input
- Gain understanding of images or video
- Mimic performance of human visual system
- Typical tasks
  - Object detection
  - Object segmentation
  - Image registration
  - Pose estimation
  - Face recognition
  - Egomotion
  - Optical Flow



# Computer Vision

- Programs that process images as input
- Gain understanding of images or video
- Mimic performance of human visual system
- Typical tasks
  - Object detection
  - Object segmentation
  - Image registration
  - Pose estimation
  - Face recognition
  - Egomotion
  - Optical Flow



# Computer Vision


- Programs that process images as input
- Gain understanding of images or video
- Mimic performance of human visual system
- Typical tasks
  - Object detection
  - Object segmentation
  - Image registration
  - Pose estimation
  - Face recognition
  - Egomotion
  - Optical Flow



# Deep Learning


Popular computer vision technique

- 2012 ImageNet Challenge significantly improved by a new method called AlexNet
  - Building on technique from 1999 (LeCun)
  - That builds on technique from 1980 (Fukushima)
- Let the computer figure out itself how to solve a problem
- Very successful in nearly all areas of computer vision
  - Defining state-of-the-art
- Prerequisites / reasons for hype
  - Lots of data for a problem
  - Fast parallel architectures (GPUs)
  - New powerful libraries






# Image Classification




- Given an image tell me what it depicts
- One of a fixed number of exclusive choices
  - Image depicts one uniquely identifiable object
  - Image may only depict a certain set of objects
- Distinguishable object choices are called classes
- Correct class of an image is called label
- A classification problem with only two classes is called binary

# Image Classification Challenges




- Object may be depicted with different acquisition techniques
- Different view angles (geometry)
- Intraclass variation
- Illumination
- Deformation
- Occlusion
- Background clutter


# Representation of images in Python



# Representation of images in Python



# Representation of images in Python





- Each picture element (pixel) is composed of three values
  - R for the red component
  - G for the green component
  - B for the blue component



# Representation of images in Python

- Each picture element (pixel) is composed of three values
  - R for the red component
  - G for the green component
  - B for the blue component





# Representation of images in Python


- Each picture element (pixel) is composed of three values
  - R for the red component
  - G for the green component
  - B for the blue component
- Images are often represented in matrix structures
  - Unclear where pixel (0,0) or (1,1) is
  - Unclear which direction is given first
- Watch your data type (OpenCV is picky)
  - uint8 [0,255] (OpenCVs favorite)
  - short [-32768, 32767]
  - float32 (for visualizing)

# Image Classification



Feature Extraction

$$\begin{bmatrix} 2 \\ 5 \\ 1 \\ 8 \end{bmatrix}$$



{cat,dog}

- Linear classifier (choice today) finds hyperplane to separate sets of points
- Transform images to point representation
  - i.e. Feature Extraction
  - Low-dimensional
  - Compact representation

# Image Classification

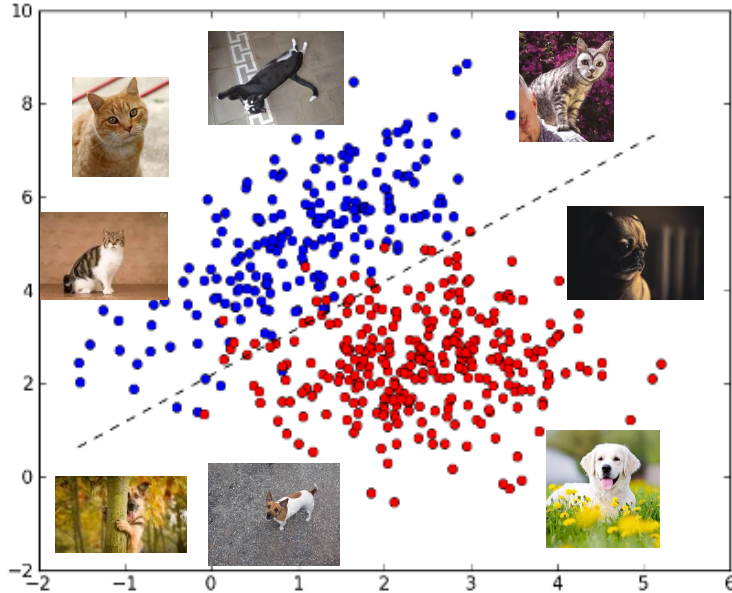
Linear classifier (choice today) finds hyperplane to separate sets of points

Transform images to point representation

- i.e. Feature Extraction
- Low-dimensional
- Compact representation


Evaluation

- Error rate:  
Percentage of wrongly classified images
- Confusion matrix:  
Error for each pair of classes



# Convolution


- Basic image processing operation: Transforms image to image
- Task: Computer similarity of each pixel with given template (kernel)





# Convolution

- Basic image processing operation
- Task: Compute similarity of each pixel with given template (kernel)
- Pay attention to range of kernel and image!



# Convolution

- Basic image processing operation
- Task: Compute similarity of each pixel with given template (kernel)
- Pay attention to range of kernel and image!

3 <sub>0</sub>	3 <sub>1</sub>	2 <sub>2</sub>	1	0
0 <sub>2</sub>	0 <sub>2</sub>	1 <sub>0</sub>	3	1
3 <sub>0</sub>	1 <sub>1</sub>	2 <sub>2</sub>	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3 <sub>0</sub>	2 <sub>1</sub>	1 <sub>2</sub>	0
0	0 <sub>2</sub>	1 <sub>2</sub>	3 <sub>0</sub>	1
3	1 <sub>0</sub>	2 <sub>1</sub>	2 <sub>2</sub>	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3	2 <sub>0</sub>	1 <sub>1</sub>	0 <sub>2</sub>
0	0	1 <sub>2</sub>	3 <sub>2</sub>	1 <sub>0</sub>
3	1	2 <sub>0</sub>	2 <sub>1</sub>	3 <sub>2</sub>
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

# Convolution

- Basic image processing operation
- Task: Computer similarity of each pixel with given template (kernel)
- Pay attention to range of kernel and image!
- Pad borders with zeros
- Stride

0 <sub>0</sub>	0 <sub>1</sub>	0 <sub>2</sub>	0	0	0	0
0 <sub>2</sub>	3 <sub>2</sub>	3 <sub>0</sub>	2	1	0	0
0 <sub>0</sub>	0 <sub>1</sub>	0 <sub>2</sub>	1	3	1	0
0	3	1	2	2	3	0
0	2	0	0	2	2	0
0	2	0	0	0	1	0
0	0	0	0	0	0	0

6.0	14.0	17.0
14.0	12.0	12.0
8.0	10.0	17.0


0	0	0 <sub>0</sub>	0 <sub>1</sub>	0 <sub>2</sub>	0	0
0	3	3 <sub>2</sub>	2 <sub>2</sub>	1 <sub>0</sub>	0	0
0	0	0 <sub>0</sub>	1 <sub>1</sub>	3 <sub>2</sub>	1	0
0	3	1	2	2	3	0
0	2	0	0	2	2	0
0	2	0	0	0	1	0
0	0	0	0	0	0	0

6.0	14.0	17.0
14.0	12.0	12.0
8.0	10.0	17.0

0	0	0	0	0 <sub>0</sub>	0 <sub>1</sub>	0 <sub>2</sub>
0	3	3	2	1 <sub>2</sub>	0 <sub>2</sub>	0 <sub>0</sub>
0	0	0	1	3 <sub>0</sub>	1 <sub>1</sub>	0 <sub>2</sub>
0	3	1	2	2	3	0
0	2	0	0	2	2	0
0	2	0	0	0	1	0
0	0	0	0	0	0	0

6.0	14.0	17.0
14.0	12.0	12.0
8.0	10.0	17.0

# Convolution



# German Traffic Sign Recognition Benchmark


- 38,000 images from (German) traffic signs
  - Vienna Convention
- 43 classes
- Over 1,000 different traffic signs instances
- Variance
  - Illumination
  - Motion Blur
  - Clutter
  - Dirt / Graffiti / Stickers
  - Occlusion
  - Angle





# German Traffic Sign Recognition Benchmark

- Filename structure
  - 0000CC/00XXX\_00YYY.ppm
    - CC = class index
    - XXX = instance of class index
    - YYY = image of instance index
- e.g. 00003/00004\_00024.ppm
  - Class 3: speed limit 60
  - Instance 4
  - Image 24
- Border of at least 5 pixel
- Border of around 10% of traffic sign size



# German Traffic Sign Recognition Benchmark

- Best human: 1.16% error rate
- Best machine classifier (2011): 0.54% error rate



# Hands-On Python: Numpy

- <https://docs.scipy.org/doc/numpy-dev/user/quickstart.html>
- “Matlab for Python”
- Matrix / Tensor manipulation (numeric)
- Fundamental library for nearly all of scientific computing in Python
- Tensorflow (Day 3 and 4) corresponds in large parts to Numpy

# Numpy: ndarray

```
import numpy as np

A = np.array( [[ 0, 1, 2], [2, 3, 4]] ) # 2x3 matrix

A.shape      # (2, 3)
A.size       # 6 (numel in Matlab)
A.dtype.name # 'int64'

B = [[0, 1, 2], [2, 3]]           # ok
A_ = np.array([[0, 1, 2], [2, 3]] ) # error
```

# Numpy: Initialization

```
A = np.array( [[ 0, 1, 2], [2, 3, 4]] ) # 2x3 matrix
```

```
B = np.zeros( (2, 3) )
```

*# 2x3 matrix*

```
C = np.ones( (2, 3, 4), dtype = np.int16 )
```

*# 2x3x4 tensor, created with data type*

```
C_ = np.zeros_like(C)
```

```
D = np.empty( (2, 3) )
```

*# 2x3 matrix, uninitialized, np.random.rand, np.random.randn*

```
E = np.arange( 0, 2, 0.3 )
```

*# array([ 0. , 0.3, 0.6, 0.9, 1.2, 1.5, 1.8])*

```
F = np.linspace( 0, 2, 9 )
```

*# array([ 0. , 0.25, 0.5 , 0.75, 1. , 1.25, 1.5 , 1.75, 2. ])*



# Numpy: Operations

*# operations usually work element-wise (even \*)*

```
a = np.array( [20,30,40,50] )
```

```
b = np.arange( 4 )
```

```
a - b           # array([20, 29, 38, 47])
```

```
b**2           # array([0, 1, 4, 9])
```

```
10*np.sin(a)    # [ 9.12945251, -9.88031624, 7.4511316 , -2.62374854]
```

```
a<35           # array([ True,  True, False, False], dtype=bool)
```

```
np.logical_and(a < 35, b > 0 )  # [False, True, False, False]
```

```
(a < 35) & (b > 0)              # brackets are necessary
```

```
a.dot(b.transpose())          # matrix product (matmul)
```

```
a.astype( np.uint8 )  # np.float32, np.int32, np.int16
```

# Numpy: Dimension manipulation

```
import numpy as np

A = np.arange( 0, 20 )

A.reshape( [4, 5] )    # 4 rows, 5 columns
A.ravel ( [4, 5] )     # back to np.arange( 0, 20 )

A.min()                # smallest element
A.min(axis = 1)        # shape = (5, 1), iterate along the columns (rowwise minimum element)
# max, sum, mean, var, cumsum

B = np.arange(0, 24).reshape( [2, 3, 4] )

C = B.sum(axis = 1)     # shape = (2, 1, 4)
C = C.squeeze()        # shape = (2, 4)

C = np.expand_dims( C, axis=1 ) # shape = (2, 1, 4)
```

# Numpy: Indexing

```
import numpy as np
```

```
A = np.arange( 0, 5 )           # [0, 1, 2, 3, 4]
```

```
A[0] = 5                       # [5, 1, 2, 3, 4]
```

```
A[2:3] = 4                     # [5, 1, 4, 3, 4]
```

```
A[-2] = 4                      # [5, 1, 4, 4, 4], A.shape[0] - 2 = 5 - 2 = 3
```

```
A[1:4:2] = 0                   # [5, 0, 4, 0, 4], start with 1, stepwidth 2, stay below 4
```

```
A[1:-1:2] = 0                 # [5, 0, 4, 0, 4], start with 1, stepwidth 2, stay below A.shape[0] - 1 = 5 - 1
```

```
A[1::2] = 0                   # [5, 0, 4, 0, 4], start with 1, stepwidth 2, stay below A.shape[0] = 5
```

# Numpy: Indexing

```
import numpy as np

A = np.arange( 0, 6 ).reshape( [2, 3] )           # [[0, 1, 2], [3, 4, 5]]

A[0, 1] = -1                                     # [[0, -1, 2], [3, 4, 5]]

A[1, :] = -1                                     # [[0, -1, 2], [-1, -1, -1]]

A[1, 1:] = 6                                     # [[0, -1, 2], [-1, 6, 6]]

A[ np.array([[True, False, True][True, True, False]], dtype=bool) ] = 1   # [[1, -1, 1], [1, 1, 6]]

A[ A > 1 ] = -1                                  # [[1, -1, 1], [1, 1, -1]]
```

# Numpy: Concatenating

```
import numpy as np

a = np.array([[1, 2], [3, 4]])           # [[1, 2], [3, 4]]
b = np.array([[5, 6]])                  # [5, 6]

np.concatenate( ( a, b ), axis = 0 )     # [[1, 2], [3, 4], [5, 6]], same as np.hstack( (a,b) )
np.concatenate( ( a, b.transpose() ), axis = 1 ) # [[1, 2, 5], [3, 4, 6]], same as np.vstack( (a,b.transpose()) )
```

# Hands-On Python: OpenCV

- see the handout for some important functions
- OpenCV can work with numpy-arrays
- But: Be careful about data types
  - Use uint8 if working in OpenCV
  - Rather receive wrong results than errors

**QUESTIONS?  
EXERCISES.**