



**RUB**

# **INTRODUCTION TO DEEP LEARNING FOR COMPUTER VISION**

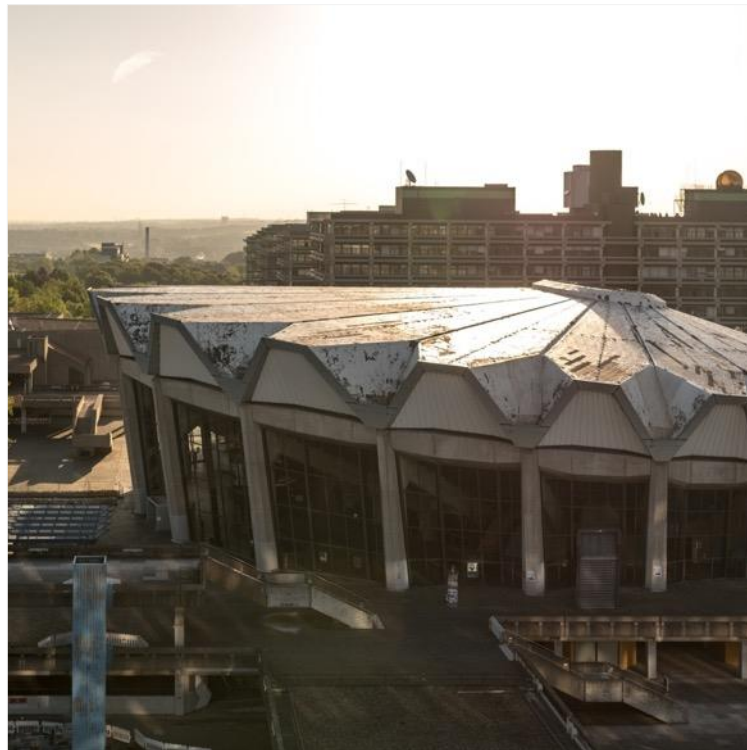
## **DAY 2 – FEATURE-BASED IMAGE CLASSIFICATION**

SEBASTIAN HOUBEN

# Schedule

## Today

- Histogram of Oriented Gradients (HOG)
- Dimensionality Reduction with Principal Component Analysis (PCA)
- Going Deeper into Classification
  - Underfitting / Overfitting
  - Training-Test-Validation
- Support Vector Machine (SVM)
- Multi-Class SVM

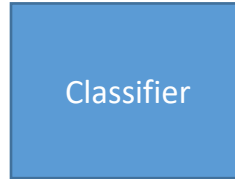


# Classification pipeline



Feature Extraction

$$\begin{bmatrix} 2 \\ 5 \\ 1 \\ 8 \end{bmatrix}$$



{cat,dog}

# Classification pipeline (Multi-class)



## Feature Extraction

Histogram-of-Oriented-Gradients

$$\begin{bmatrix} 2 \\ 5 \\ 1 \\ 8 \end{bmatrix}$$

Multi-class  
SVM

{speed limit 20,  
speed limit 30,  
..., derestriction,  
yield way, ...}

# Histogram-of-Oriented-Gradients

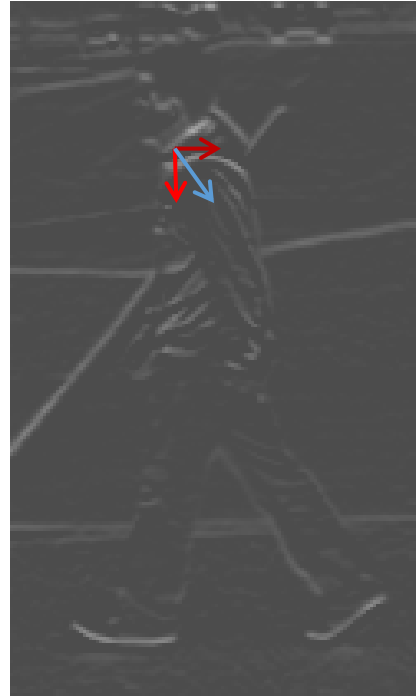
- Dalal & Triggs 2005
- Initially used for pedestrian detection
- Describes local gradient orientation distribution





# Histogram-of-Oriented-Gradients

- Dalal & Triggs 2005
- Initially used for pedestrian detection
- Describes local gradient orientation distribution
- Compute gradients
  - Convolute image with
$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} \text{ and } \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$
  - Yields pixel-wise orientation



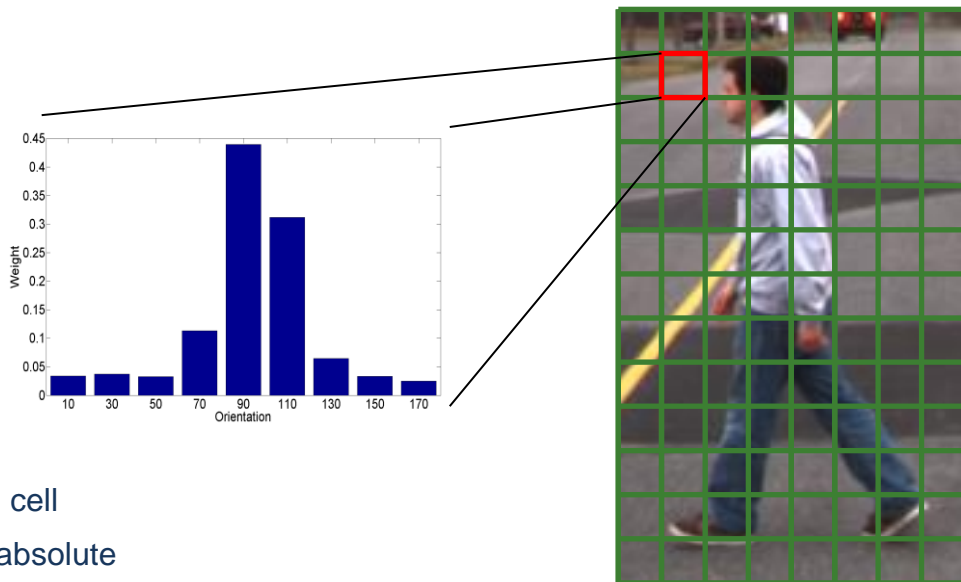
# Histogram-of-Oriented-Gradients

- Dalal & Triggs 2005
- Initially used for pedestrian detection
- Describes local gradient orientation distribution
- Compute gradients
  - Convolute image with
$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} \text{ and } \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$
  - Yields pixel-wise orientation
- Divide image into cells (e.g., 8x8 pixels)



# Histogram-of-Oriented-Gradients

- Dalal & Triggs 2005
- Initially used for pedestrian detection
- Describes local gradient orientation distribution
- Compute gradients
  - Convolute image with
$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} \text{ and } \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$
  - Yields pixel-wise orientation
- Divide image into cells (e.g., 8x8 pixels)
- Compute a histogram of all orientations present in each cell
  - Weigh the contribution of each pixel with its absolute gradient magnitude





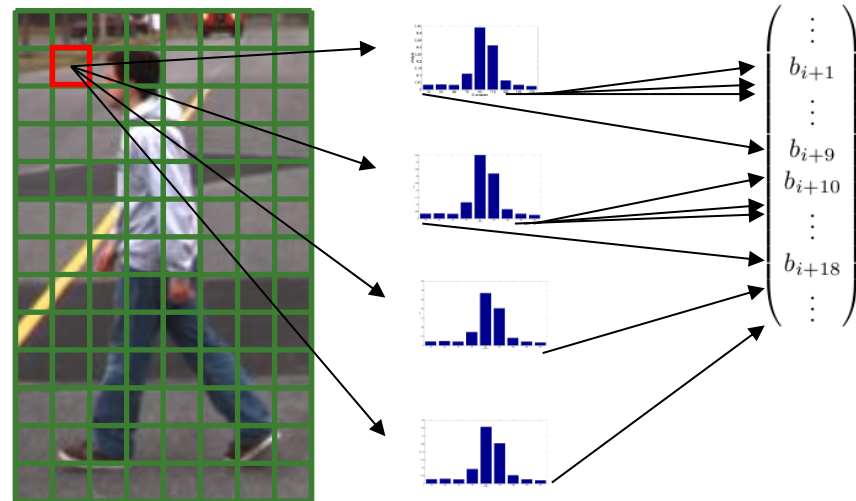
# Histogram-of-Oriented-Gradients

- Dalal & Triggs 2005
- Initially used for pedestrian detection
- Describes local gradient orientation distribution
- Compute gradients
  - Convolute image with  $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$  and  $\begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$
  - Yields pixel-wise orientation
- Divide image into cells (e.g., 8x8 pixels)
- Compute a histogram of all orientations present in each cell
  - Weigh the contribution of each pixel with its absolute gradient magnitude
- Combine neighbouring cells to blocks (e.g. 2x2 cells) and normalize histograms with respect to sum of all pixel gradients magnitudes

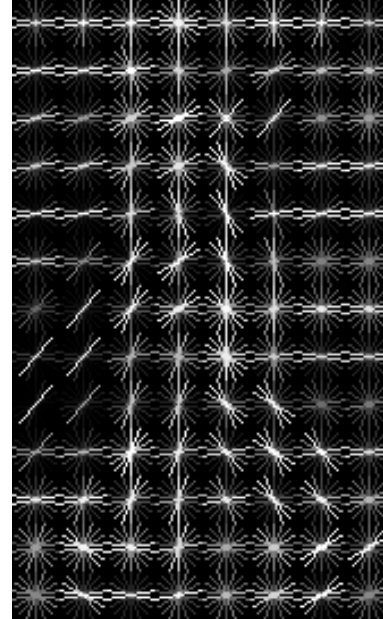


# Histogram-of-Oriented-Gradients

- Dalal & Triggs 2005
- Initially used for pedestrian detection
- Describes local gradient orientation distribution
- Compute gradients
  - Convolute image with  $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$  and  $\begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$
  - Yields pixel-wise orientation
- Divide image into cells (e.g., 8x8 pixels)
- Compute a histogram of all orientations present in each cell
  - Weigh the contribution of each pixel with its absolute gradient magnitude
- Combine neighbouring cells to blocks (e.g., 2x2 cells) and normalize histograms with respect to sum of all pixel gradients magnitudes
- For all blocks for all cells concatenate the histograms

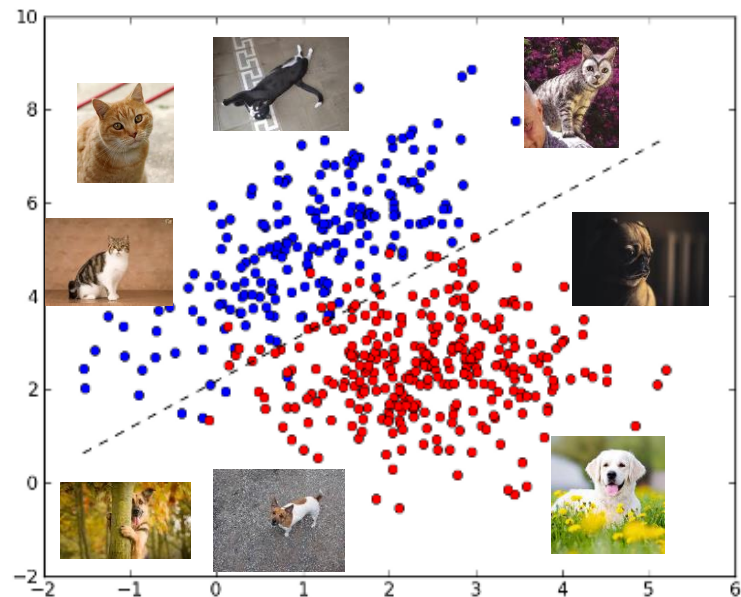


# Histogram-of-Oriented-Gradients

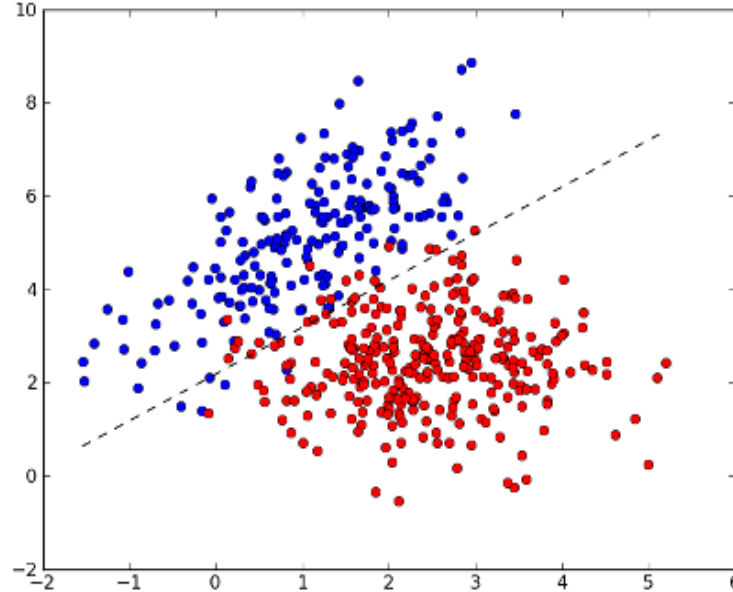


# Visualizing High-dimensional Feature Spaces

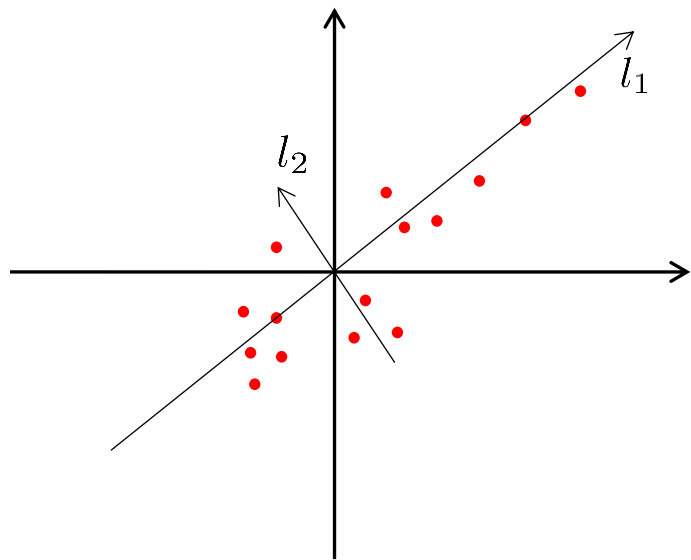
- High-dimensional vectors are hard to interpret
- Visualizing in 2d or 3d is preferable
- Dimensionality reduction / embedding
- Several methods:
  - PCA (Principal Component Analysis)
  - t-SNE (t-distributed Stochastic Nearest-Neighbour Embedding)
  - LLE (Locally-Linear Embedding)
  - MDS (Multi-Dimensional Scaling)



# Visualizing High-dimensional Feature Spaces



# Visualizing High-dimensional Feature Spaces



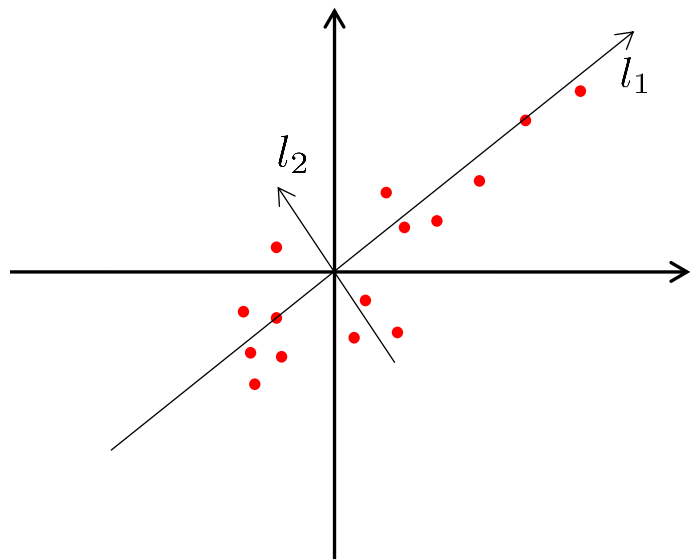
- Find function that maps data points to 2 dimensions:  $f : \mathbb{R}^n \rightarrow \mathbb{R}^2$
- Make it easy: Linear
- Thus, can be represented by a  $2 \times n$  matrix

$$f(x) = Lx = \begin{pmatrix} l_1 \\ l_2 \end{pmatrix} x$$

- But linear means 0 is mapped to 0
  - Subtract mean value from dataset beforehand
- Consists of two rows
- Rows represent the axes of main variance (principal axes)



# Visualizing High-dimensional Feature Spaces

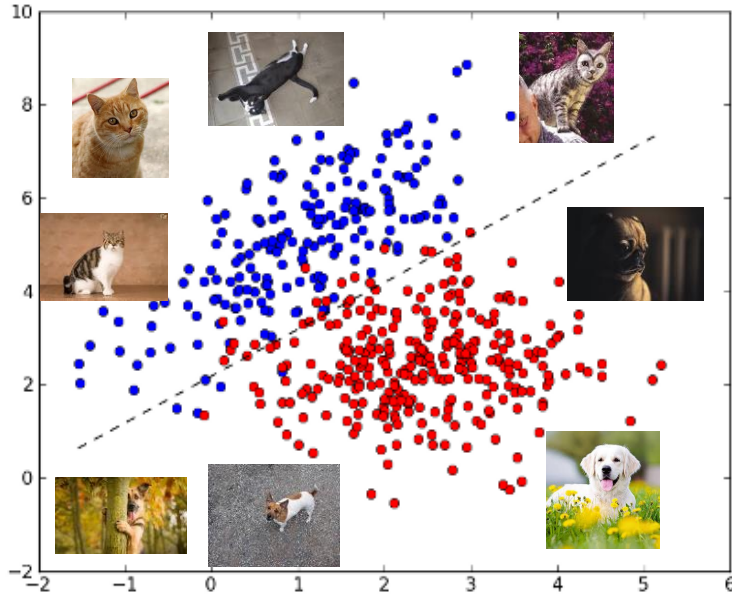


- Rows represent the axes of main variance (principal axes)
- $$\min_{l_1, \|l_1\|^2=1} \sum_i (x_i^T l_1)^2$$
- Row vector maximizing this, is given by eigenvector of 
$$C = \sum_i x_i x_i^T$$
 w.r.t. largest eigenvalue (covariance matrix  $C$ )
- Generally: Take the eigenvectors corresponding to the largest eigenvalues of the covariance matrix and project the zero-mean dataset to these vectors

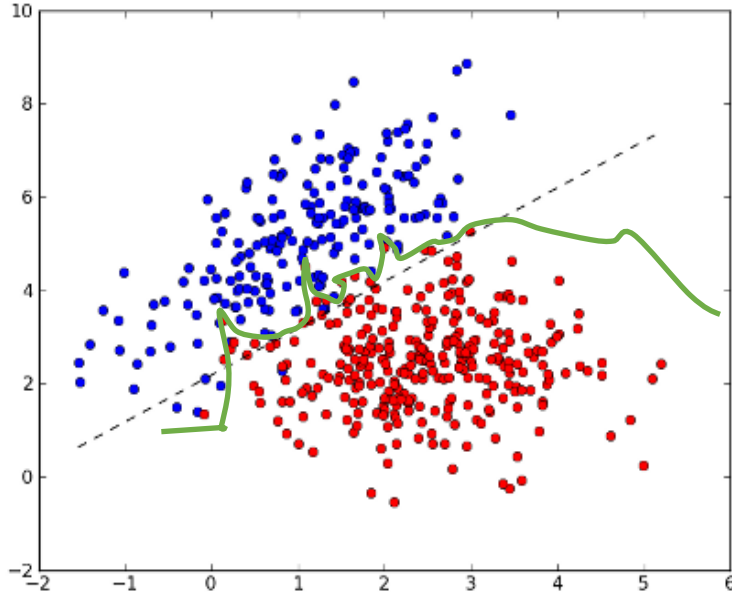
$$\begin{aligned}
& \min_{l_1, \|l_1\|^2=1} \sum_i (x_i^T l_1)^2 \\
&= \min_{l_1, \|l_1\|^2=1} \sum_i l_1^T x_i x_i^T l_1 \\
&= \min_{l_1} \max_{\lambda \geq 0} \sum_i l_1^T x_i x_i^T l_1 + \lambda (\|l_1\|^2 - 1) \\
& \min_{\lambda \geq 0} 2 \sum_i x_i x_i^T l_1 + 2\lambda l_1 = 0 \\
& \min_{\lambda \geq 0} \left( \sum_i x_i x_i^T \right) l_1 = \lambda l_1
\end{aligned}$$

# Image Classification

- Linear classifier finds hyperplane to separate sets of points
- A more complex classifier might find a better way to separate the two datasets

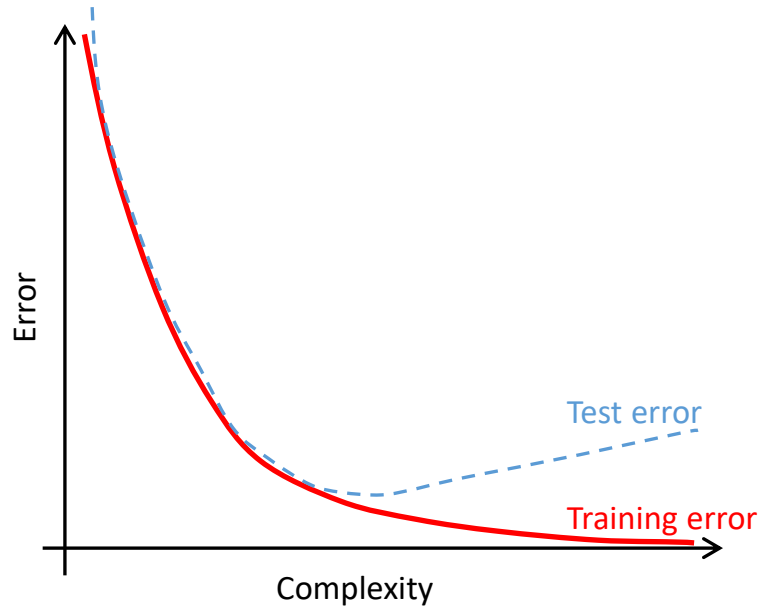


# Image Classification



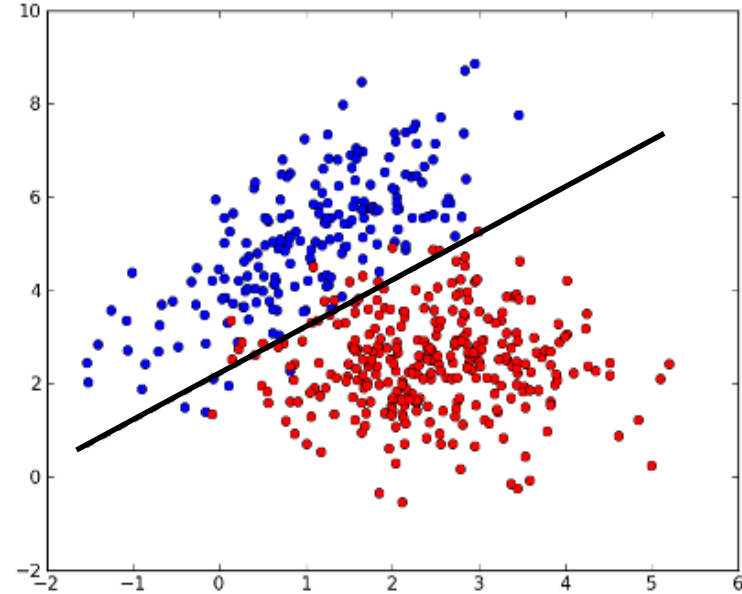
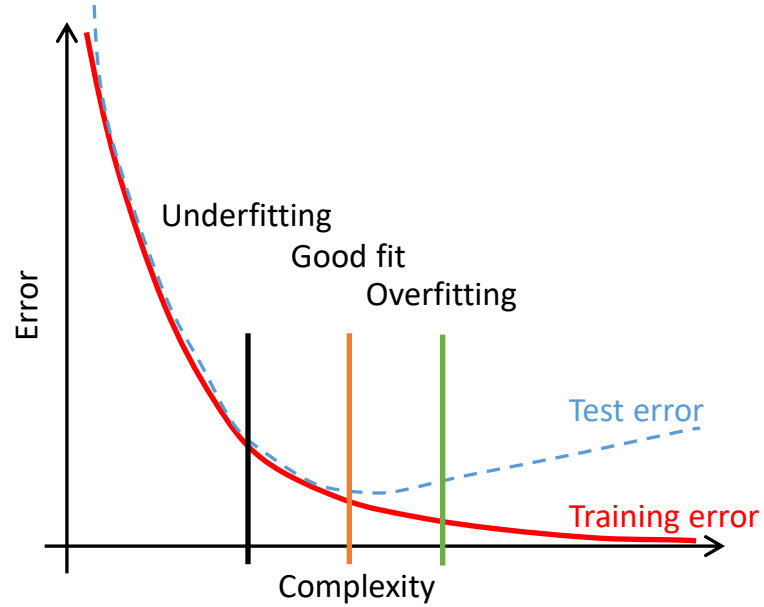
- Linear classifier finds hyperplane to separate sets of points
- A more complex classifier might find a better way to separate the two datasets
- Many ML methods have hyper-parameters that control the complexity of the function to fit
- But: In general, very complex functions tend to perform worse on unseen data

# Image Classification



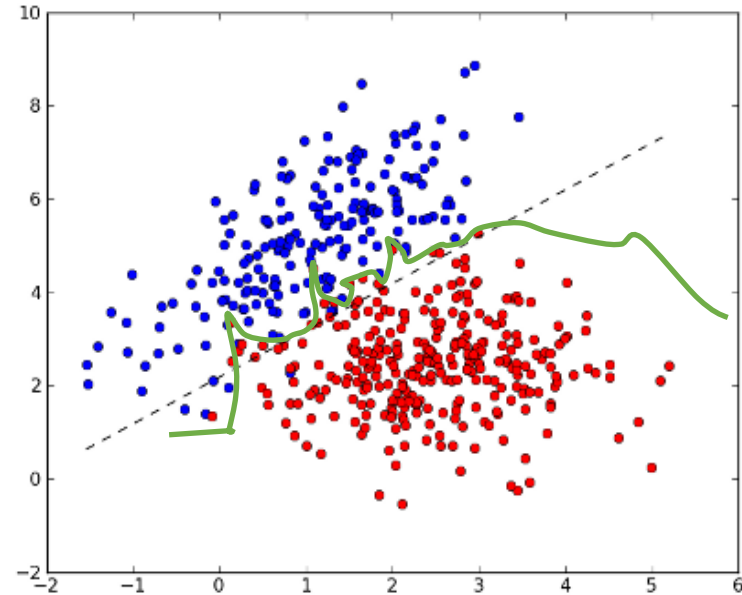
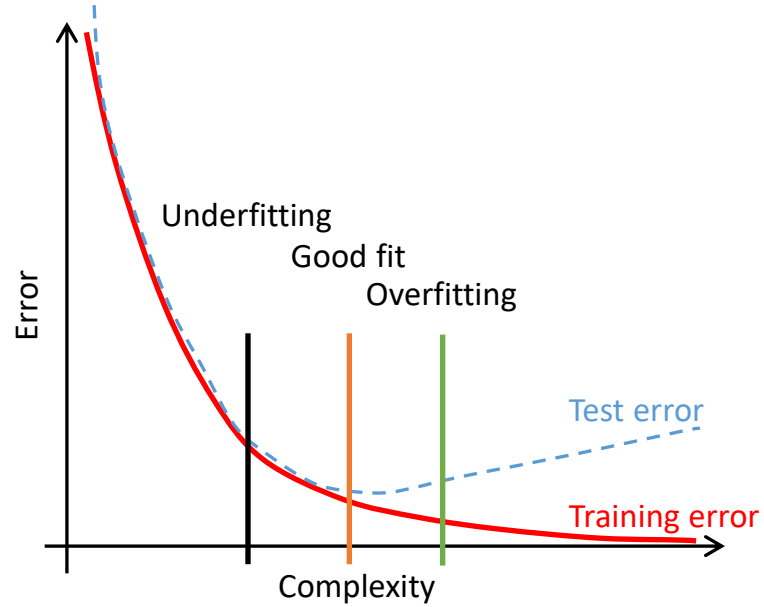
- Linear classifier finds hyperplane to separate sets of points
- A more complex classifier might find a better way to separate the two datasets
- Many ML methods have hyper-parameters that control the complexity of the function to fit
- But: In general, very complex functions tend to perform worse on unseen data

# Image Classification

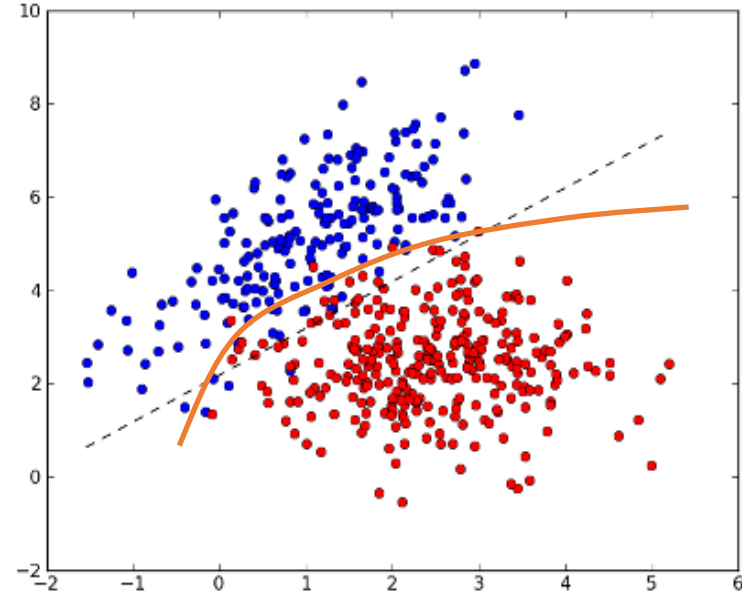
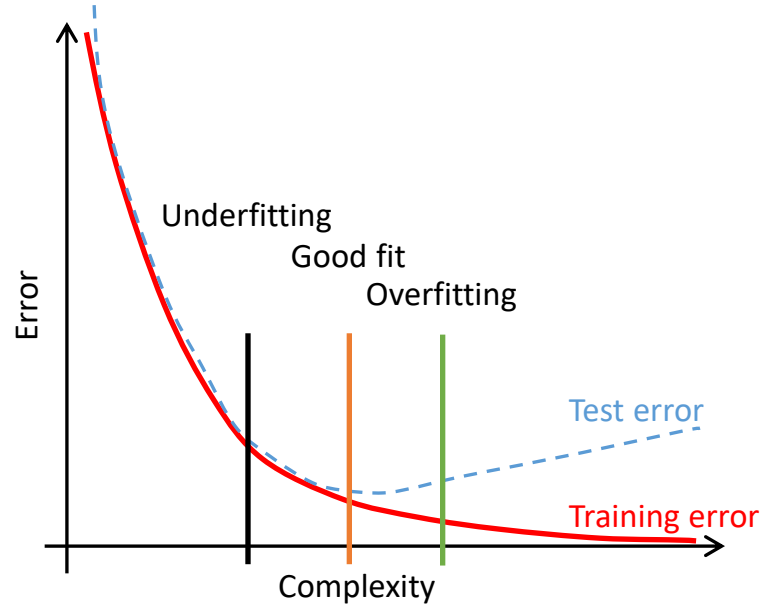




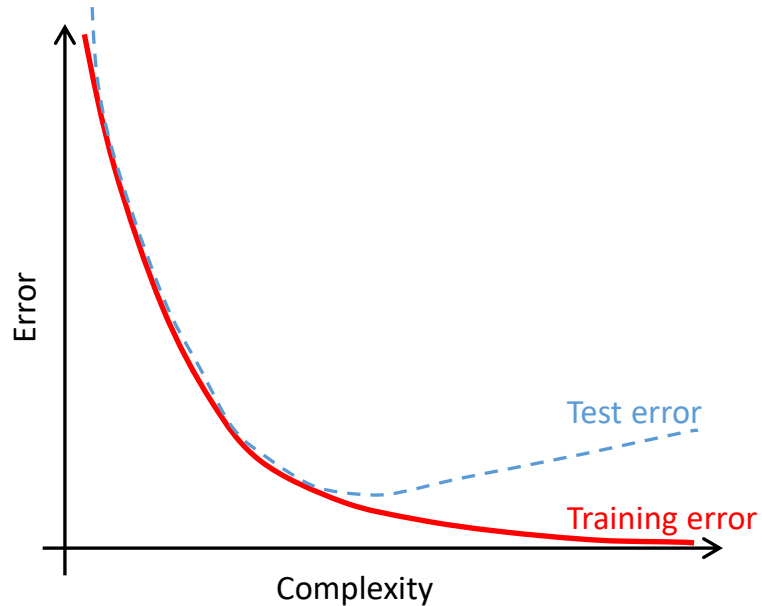
# Image Classification



# Image Classification

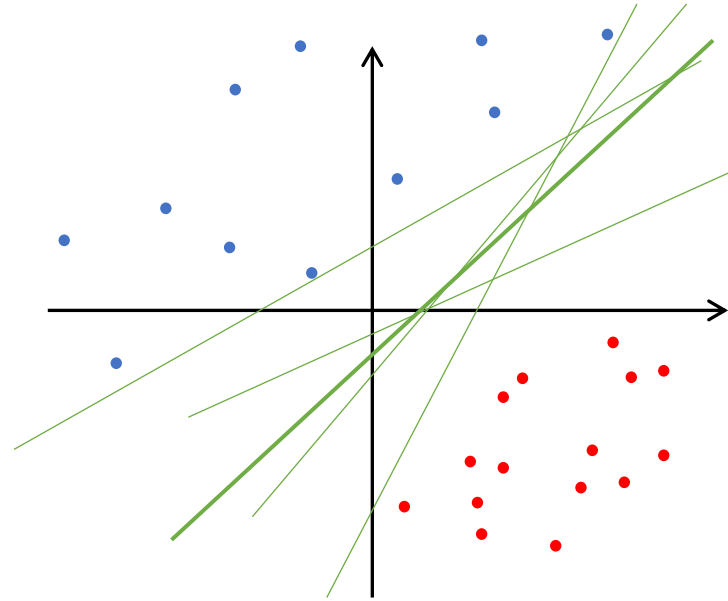


# Image Classification



- Linear classifier finds hyperplane to separate sets of points
- A more complex classifier might find a better way to separate the two datasets
- Many ML methods have hyper-parameters that control the complexity of the function to fit
- But: In general, very complex functions tend to perform worse on unseen data
- Need to estimate the training error: split dataset into training-validation-test

# Support Vector Machines



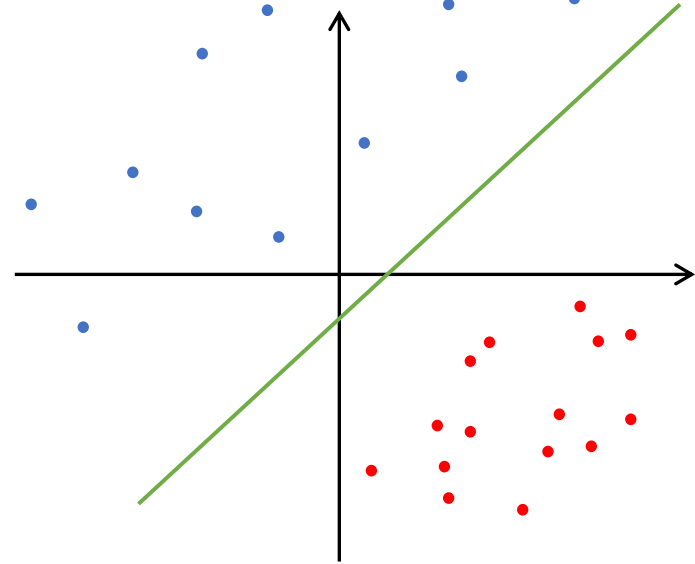
# Support Vector Machines

- Labelled Data:  $(x_i, y_i), x_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

- Solve:

$$\min_{w, b} w^T w$$

$$\text{s.t. } y_i \cdot (w^T x_i + b) \geq 1$$



# Support Vector Machines

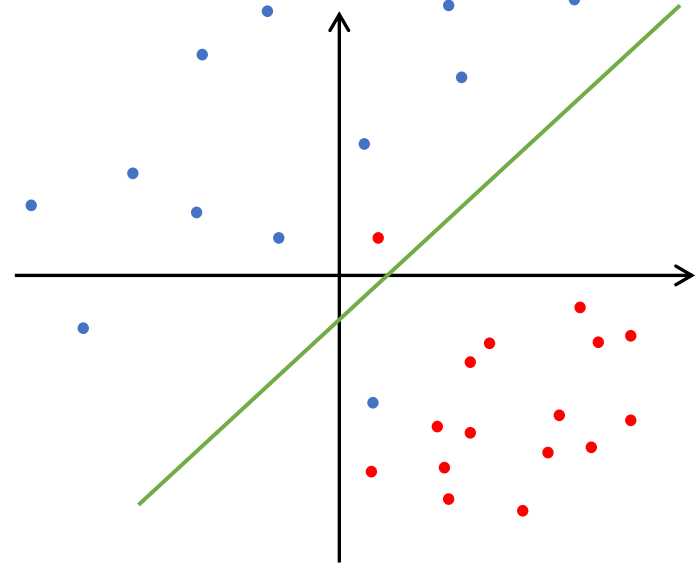
- Labelled Data:  $(x_i, y_i), x_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

- Solve:

$$\min_{w, b} w^T w + C \sum_i \xi_i$$

$$\text{s.t. } y_i \cdot (w^T x_i + b) \geq 1 - \xi_i$$

$$\xi_i \geq 0$$





# Support Vector Machines

- Labelled Data:  $(x_i, y_i), x_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

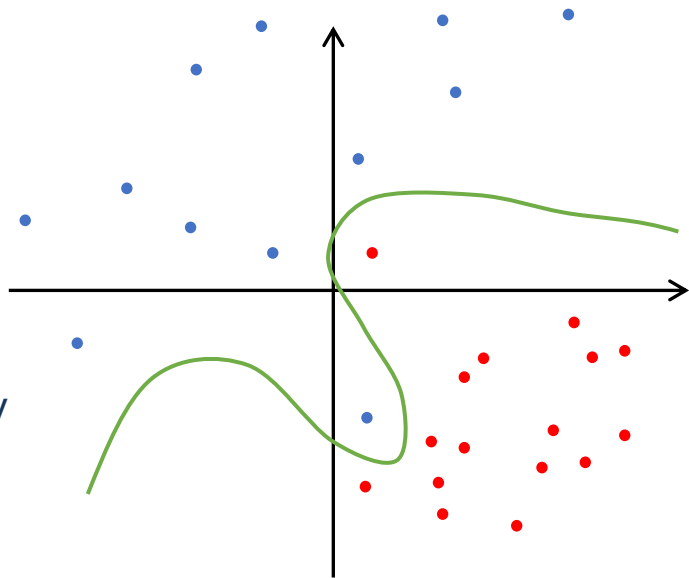
- Solve:

$$\min_{w, b} w^T w + C \sum_i \xi_i$$

$$\text{s.t. } y_i \cdot (w^T \phi_\gamma(x_i) + b) \geq 1 - \xi_i$$

$$\xi_i \geq 0$$

- $C, \gamma$  are a hyper-parameter that control complexity
- Multiclass: One-vs-All
  - most confident classifier wins
  - Confidence is given by distance to border
- Multiclass: One-vs-One



**QUESTIONS?  
EXERCISES.**